

Cycle 1: Python Basics & Mathematical Concepts

1. **What is a prime number? How can you determine if a number is prime?**
 - A prime number is a number greater than 1 that has no divisors other than 1 and itself. To check if a number is prime, check if it is divisible by any number from 2 to the square root of the number.
2. **Explain how the Fibonacci series is generated.**
 - The Fibonacci series is generated by adding the two previous numbers to get the next. Starting with 0 and 1, the series looks like: 0, 1, 1, 2, 3, 5, 8, etc.
3. **What are the types of triangles based on side lengths?**
 - Equilateral (all sides equal), Isosceles (two sides equal), Scalene (all sides unequal).
4. **Define coprime numbers.**
 - Two numbers are coprime if their greatest common divisor (GCD) is 1.
5. **How do you find the roots of a quadratic equation?**
 - Use the quadratic formula: $x = \frac{-b \pm \sqrt{b^2 - 4ac}}{2a}$, where $ax^2 + bx + c = 0$.
6. **What is a perfect number?**
 - A perfect number is a number equal to the sum of its proper divisors (excluding the number itself). Example: 6 (divisors: 1, 2, 3; $1 + 2 + 3 = 6$).
7. **Describe what an Armstrong number is.**
 - An Armstrong number (or narcissistic number) is a number that is equal to the sum of its digits raised to the power of the number of digits. Example: 153 ($1^3 + 5^3 + 3^3 = 153$).
8. **What are the differences between lists, tuples, dictionaries, and sets in Python?**
 - Lists: Ordered, mutable, allow duplicates.
 - Tuples: Ordered, immutable, allow duplicates.
 - Dictionaries: Unordered, mutable, key-value pairs.
 - Sets: Unordered, mutable, no duplicates.
9. **Explain how to add two lists element-wise.**
 - Use a loop or list comprehension, e.g., `result = [a + b for a, b in zip(list1, list2)]`.
10. **What is bubble sort?**

- A sorting algorithm where adjacent elements are compared and swapped if they are in the wrong order. Repeated until the list is sorted.

11. How can you count the vowels in a string using a dictionary?

- Loop through the string, check if each character is a vowel, and update the count in a dictionary.

12. What is the purpose of the `sum_of_digits` function in the provided code?

- It calculates the sum of all digits in a number.

13. How do you find absent digits in a mobile number?

- Compare the digits in the mobile number with the set of all digits (0-9) and find the missing ones.

14. Explain the logic used to check if a number is perfect.

- Calculate the sum of the divisors of the number. If the sum equals the number, it is perfect.

15. What does `cmath` library do in Python?

- Provides mathematical functions for complex numbers.

16. How do you ensure that input values are valid in your programs?

- Use exception handling (try-except), input validation loops, and input checks.

17. What is the significance of the swapped flag in the bubble sort algorithm?

- The swapped flag is used to optimize the bubble sort. If no swaps are made in a pass, the list is already sorted.

18. How can the efficiency of checking for prime numbers be improved?

- Instead of checking divisibility up to n , check up to \sqrt{n} .

19. How do you handle invalid input for a mobile number?

- Use validation checks to ensure the input is numeric and contains exactly 10 digits.

20. Why is it important to round the roots of a quadratic equation to two decimal places?

- To make the roots more readable and practical for real-world applications.
-

Cycle 2: NumPy

1. What is NumPy, and why is it important in Python programming?

- NumPy is a library for numerical computing in Python. It provides support for large multi-dimensional arrays and matrices, and mathematical functions to operate on them.

2. Explain the difference between a list and a NumPy array.

- A list is a general-purpose container for items, whereas a NumPy array is optimized for numerical operations, providing better performance and efficiency.

3. What are the main advantages of using NumPy arrays over Python lists?

- Faster operations, memory-efficient, support for multi-dimensional data, and a wide range of mathematical functions.

4. What is the shape of a NumPy array, and how is it determined?

- The shape is a tuple that represents the dimensions of the array (e.g., rows, columns).

5. What are the different ways to create a NumPy array?

- Using `np.array()`, `np.zeros()`, `np.ones()`, `np.arange()`, and `np.linspace()`.

6. Explain the concept of broadcasting in NumPy.

- Broadcasting allows NumPy to perform arithmetic operations on arrays of different shapes by automatically adjusting their dimensions.

7. What is the purpose of the `reshape()` function in NumPy?

- It changes the shape of an array without changing its data.

8. Describe the difference between deep copy and shallow copy in NumPy.

- A deep copy creates a new independent copy, while a shallow copy creates a reference to the original array.

9. What are universal functions (ufuncs) in NumPy?

- Functions that operate element-wise on arrays, e.g., `np.sqrt()`, `np.exp()`.

10. How can you perform matrix multiplication in NumPy?

- Use the `np.dot()` function or the `@` operator.

11. What is the difference between `np.sum()` and `np.mean()`?

- `np.sum()` returns the sum of the elements, and `np.mean()` returns the average.

12. Explain the purpose of the `axis` parameter in NumPy functions.

- The axis parameter specifies the axis along which the operation is performed (e.g., rows or columns).

13. What are masked arrays in NumPy?

- Arrays that allow you to mask certain elements (e.g., ignore them in calculations).

14. Describe how to index and slice a NumPy array.

- Indexing allows you to access individual elements, and slicing allows you to extract a sub-array.

15. What is the purpose of the np.unique() function?

- It returns the sorted unique elements of an array.

16. Explain the concept of data types in NumPy arrays.

- NumPy arrays support various data types like int, float, and complex, which are consistent within an array.

17. What is an array view in NumPy?

- A view is a new array object that looks at the same data, not a copy.

18. How does NumPy handle multidimensional arrays?

- NumPy provides tools to handle multi-dimensional arrays with indexing, slicing, and reshaping.

19. What is the purpose of the np.random module in NumPy?

- It generates random numbers and random arrays for various distributions.

20. What are some common applications of NumPy in data science?

- Data manipulation, matrix operations, statistical analysis, and machine learning algorithms.

Cycle 3: Data Visualization (Matplotlib & Seaborn)

1. What is Matplotlib, and what are its primary uses?

- Matplotlib is a plotting library in Python used to create static, interactive, and animated visualizations like line plots, bar plots, and histograms.

2. Describe the purpose of the plt.subplot() function in Matplotlib.

- plt.subplot() divides the figure into a grid and allows you to plot multiple subplots in the same figure.

3. How can you customize the style of a line plot in Matplotlib?

- You can customize the line style, color, marker, and width using parameters like `linestyle`, `color`, `marker`, and `linewidth`.

4. Explain the use of markers in Matplotlib plots.

- Markers are used to highlight data points on a plot. Common marker types include dots, squares, and triangles, set using the `marker` parameter.

5. What is the significance of the `grid()` function in plots?

- The `grid()` function adds grid lines to the plot, helping to read values more clearly.

6. How do you display a title on a plot and align it?

- Use `plt.title("Title")` to add a title. You can align it with `loc` (left, center, or right) or `pad` (distance from the plot).

7. What are histograms used for in data visualization?

- Histograms show the distribution of a dataset by grouping data points into bins or intervals.

8. Describe the role of the `bar()` function in Matplotlib.

- The `bar()` function is used to create bar charts, which are ideal for comparing categorical data.

9. How can you create a scatter plot in Matplotlib, and what is its purpose?

- A scatter plot is created using `plt.scatter()`, and it is used to visualize relationships between two variables.

10. What is the purpose of using `plt.tight_layout()`?

- `plt.tight_layout()` adjusts the spacing between subplots to prevent overlap of titles, labels, or axes.

11. Explain the difference between the `displot()` and `histplot()` functions in Seaborn.

- `displot()` is a figure-level function for univariate distributions, whereas `histplot()` is an axes-level function for plotting histograms.

12. What types of plots can you create using the `pairplot()` function in Seaborn?

- `pairplot()` creates a grid of scatter plots for all pairwise relationships in a dataset.

13. How do you display a legend in a Matplotlib plot?

- Use `plt.legend()` to add a legend to a plot, which helps identify different data series.

14. Why is it important to provide appropriate labels for axes in a plot?

- Axis labels describe the data represented and help viewers understand the context of the visualization.

15. What is the significance of the size parameter in scatter plots?

- The size parameter controls the size of the markers in a scatter plot, allowing for a third dimension of data.

16. Describe the describe() function in Pandas and its utility.

- describe() provides a summary of statistics for numerical columns in a DataFrame, such as mean, median, and standard deviation.

17. How does the read_csv() function work in Pandas?

- read_csv() is used to load data from a CSV file into a Pandas DataFrame.

18. What are the typical visualizations you can perform with categorical data?

- Common visualizations for categorical data include bar charts, pie charts, and count plots.

19. How can you improve the visual appeal of a plot?

- Improve visual appeal by customizing colors, styles, adding grid lines, and adjusting axis labels, titles, and legends.

20. Why is data visualization crucial in data analysis?

- Data visualization helps to easily communicate insights, trends, and patterns in the data, making analysis clearer and more accessible.

Cycle 4: Machine Learning

1. What is the KNN algorithm, and how does it work?

- KNN (K-Nearest Neighbors) is a classification algorithm that assigns a class based on the majority class of its K nearest neighbors in the feature space.

2. How do you choose the value of K in KNN?

- The value of K is chosen based on cross-validation or trial-and-error. A small K can lead to overfitting, and a large K can smooth out the data.

3. What is the purpose of splitting data into training and test sets?

- Splitting the data ensures that the model is trained on one set of data and tested on a separate set to assess its performance.

4. Explain the output of the classification report in KNN.

- The classification report provides metrics such as precision, recall, f1-score, and accuracy for each class in a classification problem.

5. What are some common datasets suitable for KNN classification?

- Examples include the Iris dataset, the Titanic dataset, and the MNIST dataset for image classification.

6. What is the Naive Bayes algorithm, and in what scenarios is it used?

- Naive Bayes is a probabilistic classifier based on Bayes' theorem, assuming features are independent. It is often used for text classification, spam detection, and sentiment analysis.

7. How does Gaussian Naive Bayes differ from Bernoulli Naive Bayes?

- Gaussian Naive Bayes assumes features follow a normal distribution, while Bernoulli Naive Bayes assumes binary features (0 or 1).

8. What metrics are commonly used to evaluate classification models?

- Common metrics include accuracy, precision, recall, f1-score, confusion matrix, and ROC-AUC.

9. How do you interpret the confusion matrix?

- The confusion matrix shows the counts of true positives, false positives, true negatives, and false negatives, allowing you to assess model performance.

10. What is the Decision Tree algorithm, and how does it make decisions?

- Decision Trees classify data by splitting it at each node based on feature values, creating branches leading to decision outcomes.

11. Explain how misclassification is determined in classification tasks.

- Misclassification occurs when the predicted label does not match the true label.

12. What is the importance of preprocessing data in machine learning?

- Preprocessing ensures that the data is in a suitable format, normalized, and cleaned, which improves model performance.

13. Describe the linear regression algorithm and its assumptions.

- Linear regression models the relationship between a dependent variable and one or more independent variables, assuming linearity, homoscedasticity, independence, and normality.

14. What is the difference between simple and multiple linear regression?

- Simple linear regression uses one independent variable, while multiple linear regression uses two or more independent variables.

15. How can you assess the performance of a regression model?

- Using metrics like Mean Squared Error (MSE), R-squared, and Adjusted R-squared.

16. What role does a neural network play in predicting house prices?

- A neural network can model complex non-linear relationships between features (e.g., square footage, location) and house prices.

17. Explain the significance of the activation function in a neural network.

- Activation functions introduce non-linearity, allowing the network to model complex relationships in the data.

18. How does the train-validation-test split work in neural networks?

- Data is split into three sets: training (for training the model), validation (for tuning hyperparameters), and test (for final evaluation).

19. What is the purpose of using Min-Max scaling in data preprocessing?

- Min-Max scaling normalizes the data, bringing it into a specific range, typically [0, 1], which helps improve model performance.

20. What are some common libraries used for implementing machine learning algorithms in Python?

- Common libraries include Scikit-learn, TensorFlow, Keras, and PyTorch.

Cycle 5: Web Crawling, NLP, and Clustering

1. What is a web crawler, and how does it work?

- A web crawler automatically browses the internet, extracts data from web pages, and stores it for analysis.

2. How can you extract specific HTML elements, like <p> tags, using BeautifulSoup?

- Use BeautifulSoup's `find()` or `find_all()` methods to extract tags, e.g., `soup.find_all('p')`.

3. Why is it important to check for empty links when extracting hyperlinks?

- Empty links can lead to errors or result in missing or incomplete data when scraping websites.

4. Explain the concept of n-grams in Natural Language Processing.

- N-grams are sequences of N consecutive words or characters in a text. They are useful in NLP for capturing context.

5. **How would you implement n-grams without using a library?**
 - Split the text into words, then use a loop to create sequences of N words.
6. **What are the advantages of using libraries like NLTK for NLP tasks?**
 - NLTK provides built-in functions for tokenization, stemming, stop word removal, and many other NLP tasks, saving time and effort.
7. **What is the purpose of stop word removal in text processing?**
 - Stop word removal eliminates common but irrelevant words (e.g., "the," "is") to focus on meaningful words in text analysis.
8. **Describe the process of tokenization in NLP.**
 - Tokenization is the process of splitting text into smaller units, such as words or sentences, for further analysis.
9. **How does K-means clustering work?**
 - K-means is an iterative algorithm that divides data into K clusters by minimizing the variance within each cluster.
10. **What factors should you consider when choosing the number of clusters (K) in K-means?**
 - Factors include the dataset's size, the data's natural grouping, and the "elbow method" to find an optimal K.
11. **How can you visualize clusters in a scatter plot?**
 - Use colors or markers to distinguish different clusters in a scatter plot.
12. **Why is it important to visualize the cluster centers in K-means?**
 - Cluster centers represent the average of all points in a cluster and provide insight into the structure of the data.
13. **What does it mean to display cluster labels for each point?**
 - Cluster labels assign each data point to its respective cluster, making it easier to analyze groupings.
14. **How can the elbow method help in selecting K for K-means?**
 - The elbow method involves plotting the sum of squared distances (within-cluster variance) for different values of K. The "elbow" point indicates the optimal K.
15. **What are the potential challenges when implementing K-means clustering?**
 - Challenges include selecting the right K, dealing with outliers, and ensuring clusters are well-separated.

16. How can different values of K affect the results of clustering?

- A low K can result in underfitting, while a high K may lead to overfitting and less meaningful clusters.

17. What are some common metrics to evaluate clustering performance?

- Metrics include Silhouette score, Davies-Bouldin index, and Inertia.

18. What are the practical applications of K-means clustering?

- Applications include customer segmentation, image compression, and anomaly detection.

19. Why is Python a popular choice for implementing web crawlers and NLP tasks?

- Python has libraries like BeautifulSoup, Scrapy, NLTK, and SpaCy, which simplify web crawling and NLP tasks.

20. How can data preprocessing impact the results of clustering algorithms like K-means?

- Proper preprocessing (e.g., scaling, handling missing values) can improve the quality of clusters and the algorithm's performance.